


# AUTOENCODERS

# A general principle of generation

- Data is encoded into a different representation
- New data is generated by sampling from the new representation
- GMMs are just one type of encoding-decoding scheme



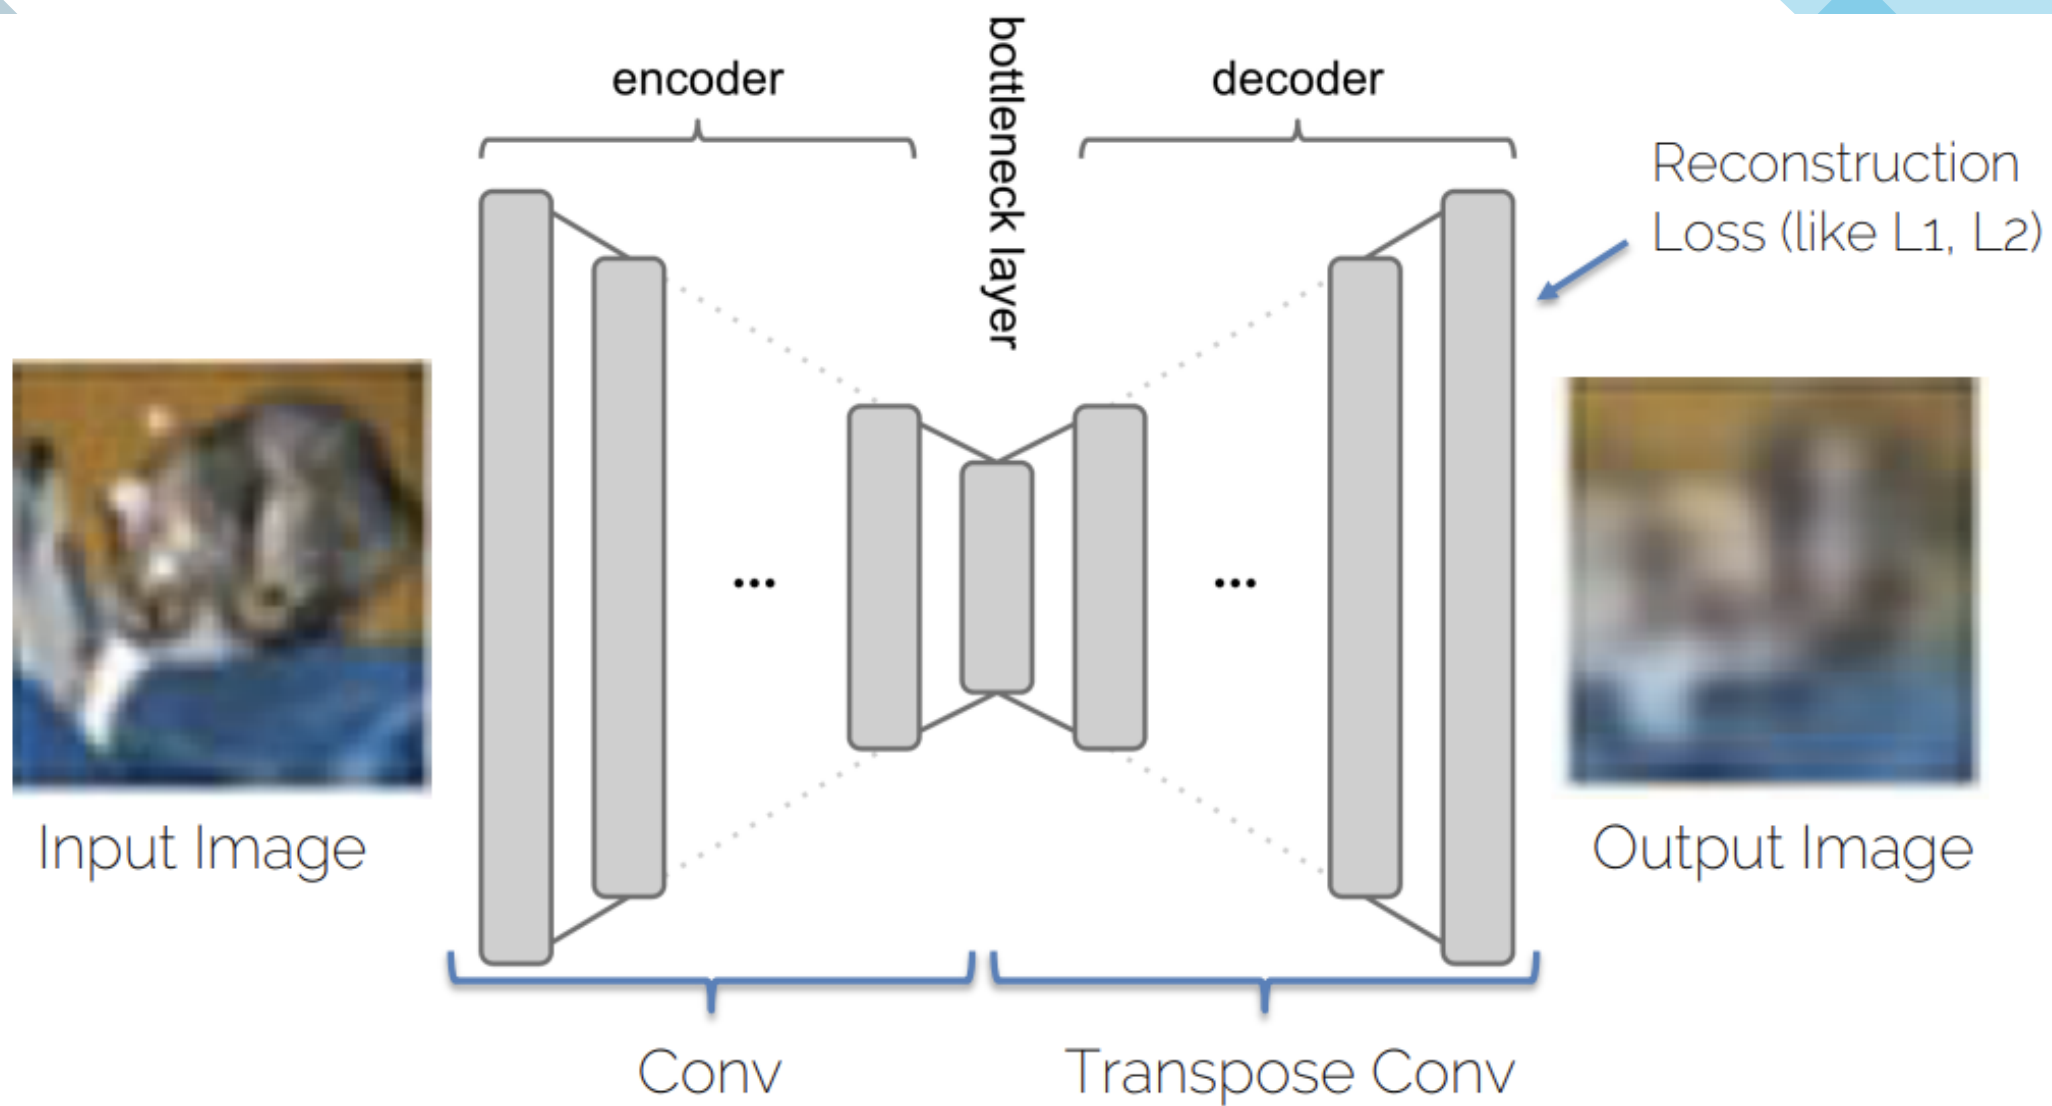
Autoencoders consist of two main parts:




**The encoder:** Compresses the input data into a lower-dimensional latent space.



**The decoder:** Reconstructs the original data from this compressed representation.






- Encoder stage : map the input  $x$  to  $z$   

$$z = \delta(Wx + b)$$
 $\delta$  : element-wise activation function
- Decoder stage : map  $z$  to reconstruction  $x'$


Autoencoders are trained to minimise :

$$L(x, x') = \|x - x'\|^2$$



# Autoencoders: applications

- Denoising: input clean image + noise and train to reproduce the clean image.



# Autoencoders: Applications

- Image colorization: input black and white and train to produce color images



# Autoencoders: Applications

- Watermark removal





# Simple bottleneck layer in Keras

- `input_img = Input(shape=(784,))`
- `encoding_dim = 32`
- `encoded = Dense(encoding_dim, activation='relu')(input_img)`
- `decoded = Dense(784, activation='sigmoid')(encoded)`
- `autoencoder = Model(input_img, decoded)`
- Maps 28x28 images into a 32 dimensional vector.




# Denoising autoencoders

- Basic autoencoder trains to minimize the loss between  $x$  and the reconstruction  $g(f(x))$ .
- Denoising autoencoders train to minimize the loss between  $x$  and  $g(f(x+w))$ , where  $w$  is random noise.
- Same possible architectures, different training data.
- [Kaggle has a dataset on damaged documents.](#)



# Denoising autoencoders

- Denoising autoencoders can't simply memorize the input output relationship.
- Intuitively, a denoising autoencoder learns a projection from a neighborhood of our training data back onto the training data.



# Sparse autoencoders

- Construct a loss function to penalize *activations* within a layer.
- Usually regularize the *weights* of a network, not the activations.

## How It Works:

- The autoencoder is trained with an additional **sparsity penalty** (like KL-divergence or L1 regularization) that forces most neurons in the hidden layer to remain inactive for most inputs.
- This encourages the network to learn a **disentangled and meaningful** representation of the input data.

# Mathematical Formulation of Sparse autoencoders

Given the activation of the hidden unit  $j$  as  $a_j$ , the sparsity constraint is applied using KL-divergence:

$$KL(\rho || \hat{\rho}) = \sum_j \rho \log \frac{\rho}{\hat{\rho}_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j}$$

where:

- $\rho$  is the target sparsity (small value, e.g., 0.05),
- $\hat{\rho}_j$  is the average activation of neuron  $j$ .

## Regularization Used:

- **L1 Regularization (Lasso):** Adds an absolute value penalty to the activations.
- **KL-Divergence:** Ensures that the average activation of hidden neurons stays close to a low desired value.


# Contractive autoencoders

## How It Works:

- Instead of enforcing sparsity, a **contractive penalty** is added, which minimizes the **sensitivity of the hidden representation to small input variations**.
- This prevents overfitting and makes the autoencoder robust to noise.

$$\mathcal{L}(x, \hat{x}) + \lambda \sum_i \left\| \nabla_x a_i^{(h)}(x) \right\|^2$$

Similar **inputs** are contracted to a constant **output** within a neighborhood, based on what the model **observed** during training



# Mathematical Formulation of Contractive Autoencoders

The contractive loss adds a penalty term:

$$\mathcal{L} = \|x - \hat{x}\|^2 + \lambda \sum_j \|\nabla_x h_j\|^2$$

- **Regularization Used:**


**Jacobian Regularization:** Penalizes large gradients of the encoder output with respect to the input.

where:

- $x$  is the input,
- $\hat{x}$  is the reconstructed output,
- $h_j$  is the activation of neuron  $j$ ,
- $\nabla_x h_j$  is the Jacobian matrix of the activations,
- $\lambda$  controls the weight of the contractive penalty.

# Contractive autoencoders

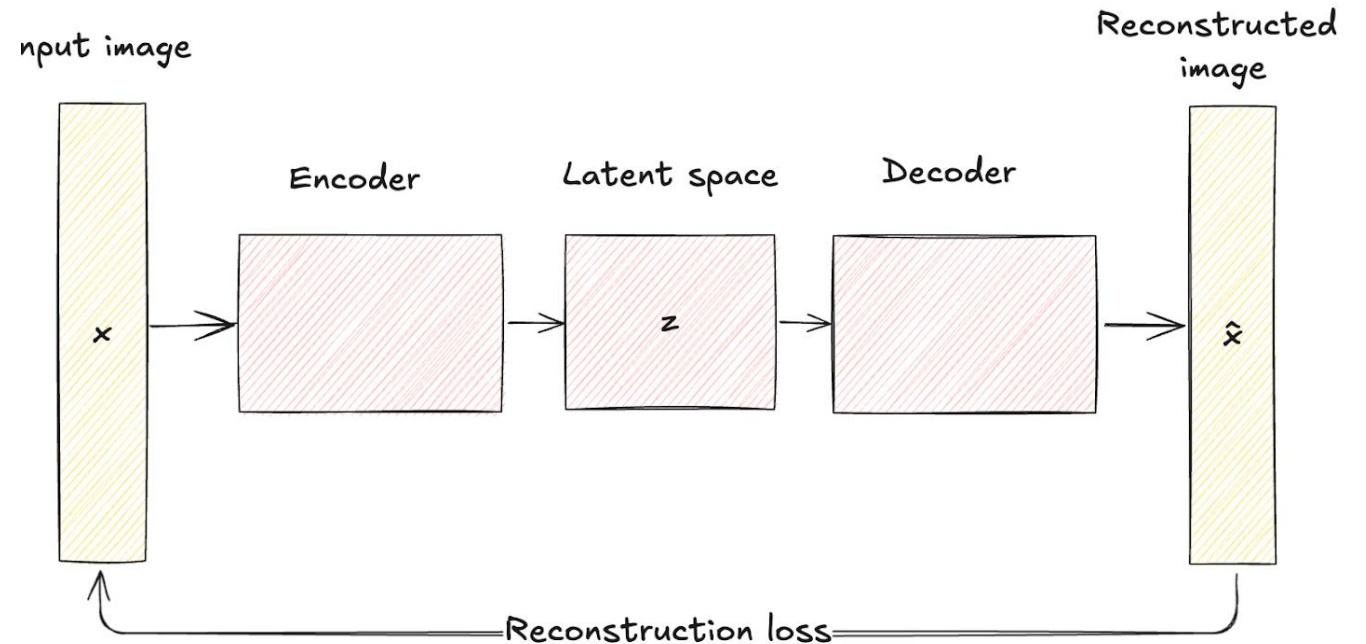
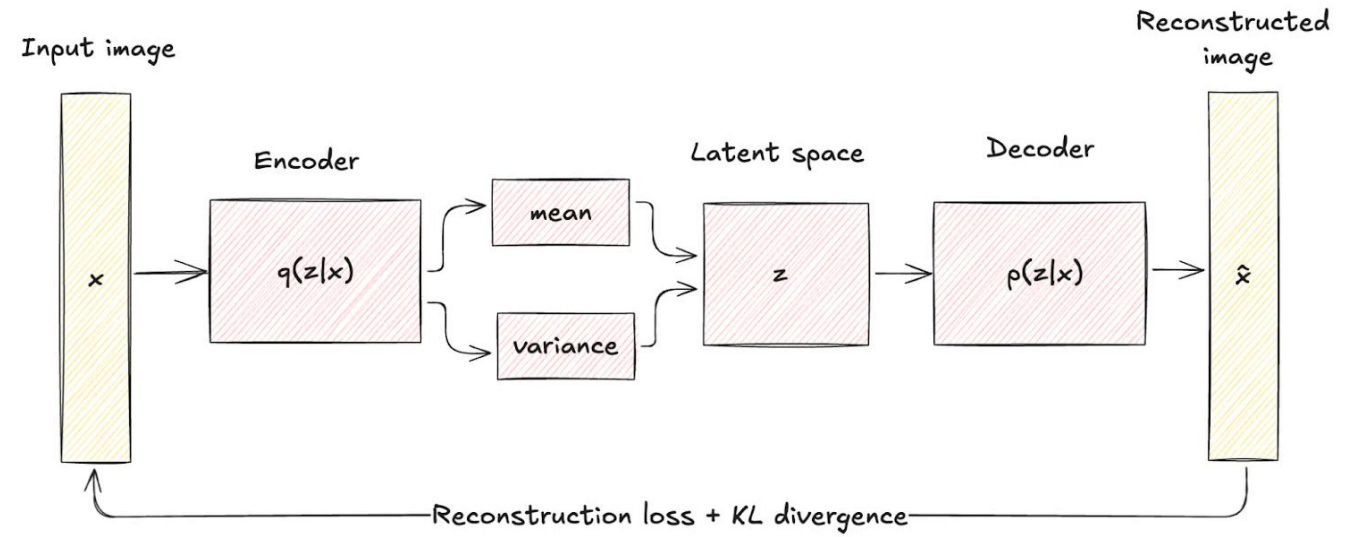
- Contractive autoencoders make the *feature extraction function* (ie. encoder) resist infinitesimal perturbations of the input.







# Variational Autoencoder

- Key innovation of VAEs lies in their ability to generate new, high-quality data by learning a structured, continuous latent space.
- Traditional autoencoders produce a fixed point in the latent space (the network that maps the input data  $x$  to a fixed, lower-dimensional latent space representation  $z$ . This process is deterministic, meaning each input is encoded into a specific point in the latent space). The decoder network then reconstructs the original data from this fixed latent representation, aiming to minimize the difference between the input and its reconstruction.
- Encoder in a VAE outputs parameters of a probability distribution—typically the mean and variance of a Gaussian distribution (encoder in a VAE maps the input data to a probability distribution over the latent variables, typically modeled as a Gaussian distribution with mean  $\mu$  and variance  $\sigma^2$ ). During training, we sample a point from this distribution to feed into the decoder.




# VAEs concept

$$z \sim N(0, 1)$$
$$x|z \sim N(\mu(z), \sigma(z))$$



We approximate complex  $p(x)$  as a composition of many simple  $z$





This approach encodes each input into a distribution rather than a single point, adding a layer of variability and uncertainty.

Point in data space



Latent space



**Figure 7:** Decoupling attribute vectors for smiling (x-axis) and mouth open (y-axis) allows for more flexible latent space transformations. Input shown at left with reconstruction adjacent. (model: VAE from Lamb 16 on CelebA)




**Figure 4.4:** VAEs can be used for image resynthesis. In this example by White, 2016, an original image (left) is modified in a latent space in the direction of a *smile vector*, producing a range of versions of the original, from smiling to sadness.



# VAE outputs



Samples from a VAE trained on MNIST



Samples from a VAE trained on a faces dataset

# VAE limitations

- People have mostly moved on from VAEs to use GANs for generating synthetic high-dimensional data
- VAEs are theoretically complex
- Don't generalize very well
- Are pragmatically under-constrained
  - Reconstruction error need not be exactly correlated with realism



Realistic



Fake



Who is real ?



[Play](#)

<https://www.whichfaceisreal.com/index.php>